TensorFlow

Adversarial Regularization

George Yu

A Quick Refresher

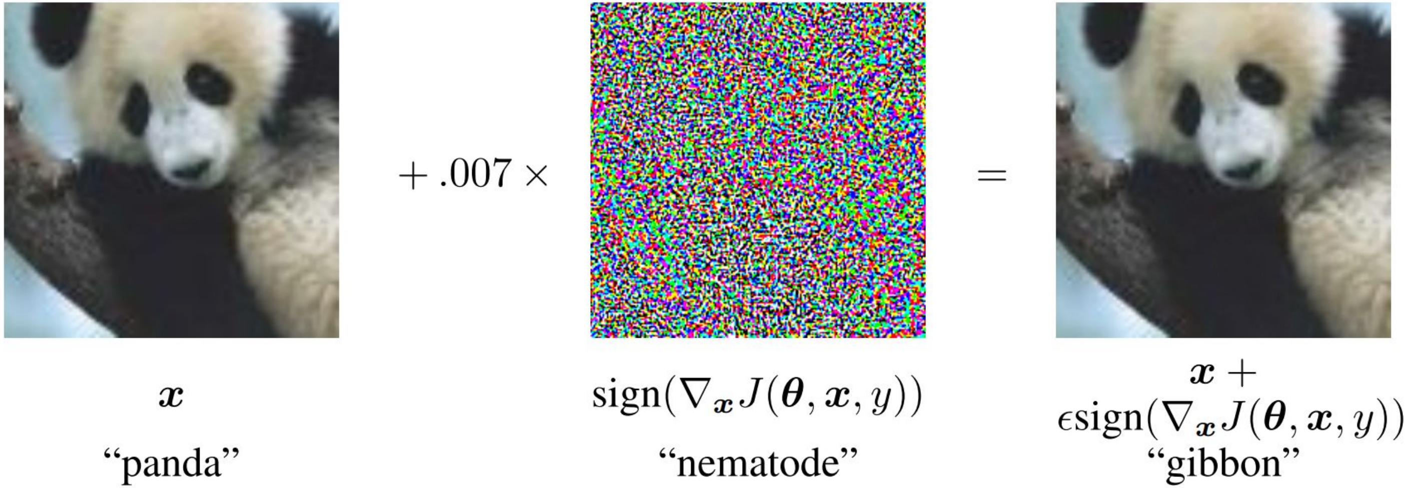
$+ .007 \times$
$=$
$\boldsymbol{x}$
"panda"
$\text{sign}(\nabla_{\boldsymbol{x}} J(\boldsymbol{\theta}, \boldsymbol{x}, y))$
"nematode"
$\boldsymbol{x} + \epsilon\text{sign}(\nabla_{\boldsymbol{x}} J(\boldsymbol{\theta}, \boldsymbol{x}, y))$
"gibbon"

[Source: Goodfellow, et al., ICLR'15]

Adversaries - FGSM

- FGSM = Fast Gradient Sign Method
- Find the direction that increases the loss the most, and take a step
- Significantly degrades accuracy of unregularized model

Adversaries - PGD

- PGD = Projected Gradient Descent
- Iteratively apply the FGSM algorithm
- Project back to within the perturbation budget after every step
- Produces a stronger adversary but is more expensive than FGSM.

Adversarial Regularization in NSL

- Generate adversarial examples as graph neighbors.
- Regularize the loss by including the neighbors in the graph.
- Colab notebook